

We now consider a model depicting a cloud-computing architecture shown in Fig (1). There is a solid, directed edge from component type i to component type j if $j \in \Gamma_i$, and the label on the edge is $\phi_{i,j}$. The dashed, undirected edge shows the connection of our system to the public internet. Our system is designed to handle hypertext transfer protocol (HTTP) requests and simple mail transfer protocol (SMTP) requests. NS represents a network switch; LB1 and LB2 are two types of load balancers; FW1 and FW2 denote two types of firewalls; HV1 and HV2 signify two types of hypervisors; and SR1 and SR2 portray two types of server racks. This organization of components specifically shows a firewall “sandwich” as described in [1].

We set $r_{NS} = r_{LB1} = r_{LB2} = r_{FW1} = r_{FW2} = r_{HV1} = r_{HV2} = 2$ and $r_{SR1} = r_{SR2} = 3$. Let v_i define the number of components of type i that need to be up for the system to remain operational. We need at least two server racks of each type SR1 and SR2 to be up to be able to parallel process across server racks. For other component types, we require at least one component of each type of the system to be up in order for the system to remain operational. Thus, $v_{NS} = v_{LB1} = v_{LB2} = v_{FW1} = v_{FW2} = v_{HV1} = v_{HV2} = 1$ and $v_{SR1} = v_{SR2} = 2$. This gives us a state space of size 69984. Solving a model of this size using the previous version of DECaf takes xxx seconds, whereas using the new version’s tree-generation algorithms given in Section ?? takes xxx seconds.

We build the system such that all component types that are located to the left of the network switch (i.e., types with a “1” in their names) handle SMTP requests only, whereas all component types located to the right of the network switch (i.e., types with a “2” in their names) handle HTTP requests only. Thus, both the branches need to be operational for the system to be able to handle HTTP and SMTP requests. This type of decoupling ensures that hardware failure on one side does not immediately propagate to the other. In this model the system up condition is an AND of the individual-up conditions. Note that both the previous and new versions of DECaf allow for a general boolean combination of inequalities.

We assume that failure of any component causes a downward-propagating cascading failure because components on levels below the failed component end up going offline and reach a dormant state. Note that the propagation cannot continue via firewalls onto server racks. This is because although the failure of a firewall renders server racks vulnerable, they are still “online”. To comply with our model, dormancy is treated as a failure. We define the public internet to be at *depth* 0, the network switch to be at *depth* = 1 and

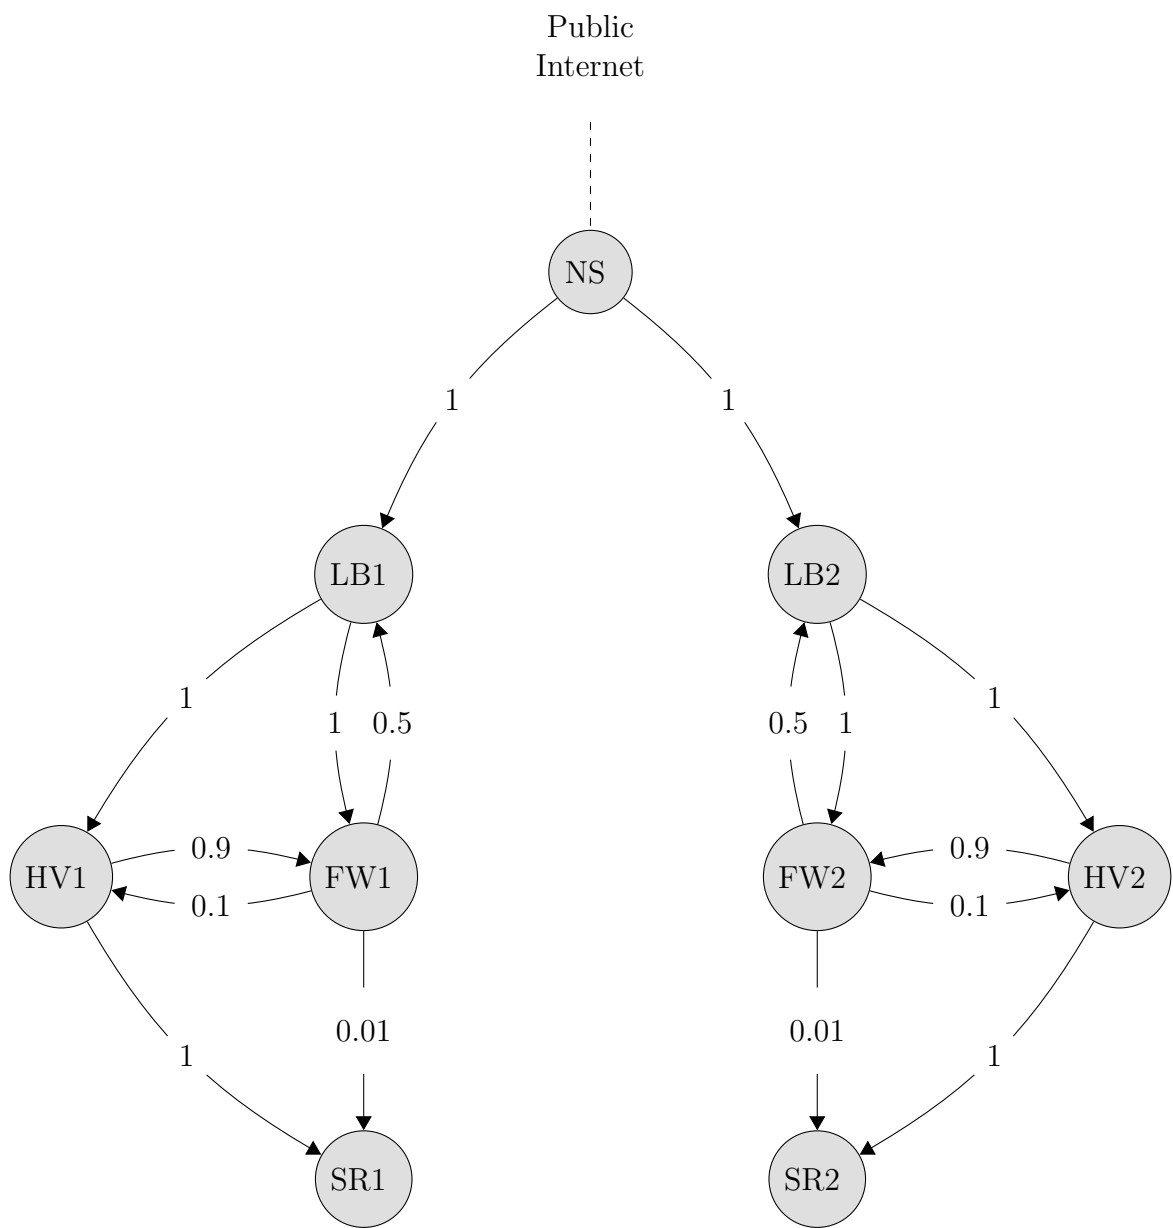


Figure 1: Cloud-computing architecture

so on. Thus for the downward-propagating cascading failure, $\phi_{i,j} = 1 \mid i, j \in \Omega, \text{depth}(j) = \text{depth}(i) + 1$ AND i and j are on the same side. As stated in [2] “hypervisors almost always cause other system components to fail” and certainly cause server racks to fail because of state corruption. Thus, there is probability 0.9 that a hypervisor causes a firewall to fail. A firewall that solely implements an (access- control-list) verification is one example of a vulnerable firewall. A firewall that is not robust may let malicious packets get through and cause a failure of a load balancer. We model the probability of this to be 0.5 assuming half the packets received are malicious. A firewall can cause a hypervisor to fail by sending malformed packets due to software bugs or hardware failure. We model the probability of this is 0.1. A firewall in rare cases (which we model by setting $\phi_{FW1,SR1} = \phi_{FW2,SR2} = 0.01$) can expose server racks to attack, though most likely server racks become inaccessible because of a hypervisor failure.

To estimate λ_i for each i , we assume that in a high-demand environment a component is going to fail on average in twice the time of its warranty and in a low-demand environment a component fails in about ten times the period of its warranty. Note that our time unit is in hours. Typical commercial network switches, load balancers and hypervisors have warranties of 90 days (2160 hours). Thus, $\lambda_{NS,0} = \lambda_{LB1,0} = \lambda_{LB2,0} = \lambda_{HV1,0} = \lambda_{HV2,0} = 1/4320$ and $\lambda_{NS,1} = \lambda_{LB1,1} = \lambda_{LB2,1} = \lambda_{HV1,1} = \lambda_{HV2,1} = 1/21600$. Since firewalls almost never fail in isolation, we set $\lambda_{FW1,0} = \lambda_{FW2,0} = \lambda_{FW2,1} = \lambda_{FW1,1} = 1 \times 10^{-10}$. Commercial server racks often have 3-year (26280 hour) warranties giving us $\lambda_{SR1,0} = \lambda_{SR2,0} = 1/52560$ and $\lambda_{SR1,1} = \lambda_{SR2,1} = 1/262800$.

For the components’ repair rates, we assume that all hardware components are swapped out with a repair rate of 1, i.e., we can replace one component an hour on average. A repair for a hardware component may also be having to restart that part of the system if the component is in a dormant condition. For simplicity we assume this rate is 1 as well. Firewalls being non-hardware components need to be reconfigured after they fail and this occurs with a repair rate of 0.1.

Additionally, our system operates in two environments: high demand ($e = 0$) and low demand ($e = 1$). Our system switches between environments once every 12 hours giving us the following environment transition rates $\nu_{0,1} = 1/12$ and $\nu_{1,0} = 1/12$.

References

- [1] Salchow Ken Jr. Load balancing 101: Firewall sandwiches. Technical report, F5 Networks, Inc, 2010.
- [2] Ye Michael and Tamir Yuval. Rehype: Enabling vm survival across hypervisor failures. 2011.