Working Title

M. Sanghavi, S. Tadepalli, M. Nakayama

January 2013

1 Model

We work with the stochastic model of [1], which considers the evolution over time of a repairable dependability system operating in a randomly changing environment. The system consists of a collection $\Omega = \{1, 2, \ldots, N\}$ of $N < \infty$ component types. Each component type $i \in \Omega$ has a redundancy $1 \leqslant r_i < \infty$, and the r_i components of type i are assumed to be identical. A component can be either operational (up) or failed (down).

The environment changes randomly within a set $\mathcal{E} = \{0, 1, 2, \ldots, L\}$. For example, we can think of the environment as representing the current load on the system, and if there are two possible environments, 0 and 1, then 0 (resp., 1) may represent a low (resp., high) load. Once the environment enters $e \in \mathcal{E}$, it remains there for an exponentially distributed amount of time with rate $\nu_e > 0$, after which the environment changes to e' with probability $\delta_{e,e'} \geqslant 0$, where $\delta_{e,e} = 0$ and $\sum_{e' \in \mathcal{E}} \delta_{e,e'} = 1$. We assume the matrix $\delta = (\delta_{e,e'} : e, e' \in \mathcal{E})$ is irreducible; i.e., for each $e, e' \in \mathcal{E}$, there exists $k \geqslant 1$ and a sequence $e_0 = e, e_1, e_2, \ldots, e_k = e'$ with each $e_i \in \mathcal{E}$ such that $\prod_{i=0}^{k-1} \delta_{e_i, e_{i+1}} > 0$. In other words, it is possible to eventually move from environment e to environment e'.

The components in the system can randomly fail and then be repaired. When the environment is $e \in \mathcal{E}$, the failure rate and repair rate of each component of type i are $\lambda_{i,e} > 0$ and $\mu_{i,e} > 0$, respectively. If there is only one environment e, i.e., $|\mathcal{E}| = 1$, then the lifetimes and repair times of components of type i are exponentially distributed with rates $\lambda_{i,e}$ and $\mu_{i,e}$, respectively. Exponential distributions are frequently used to model lifetimes of hardware and software components; e.g., see [2]. We assume that all operating components of a type i have the same failure rate $\lambda_{i,e}$ in environment e. Thus, in a system with redundancies for which not all components of a type are needed for operation of the system, the extras are "hot spares" since they fail at the same rate as the main components.

Our model includes probabilistic instantaneous cascading failures, which occur as follows. The ordered set Γ_i specifies the types of components that a failure of a type-i component can cause to simultaneously fail. When a component of type i fails, it causes a single component of type $j \in \Gamma_i$ to fail at the same time with probability $\phi_{i,j}$ (if there are components of type j up) and we call $\phi_{i,j}$ a *component-affected probability*. The events that the individual components of types $j \in \Gamma_i$ fail immediately are independent. Thus, when a component of type i fails, there are independent "coin flips" to determine which components in Γ_i fail, where the coin flip for $j \in \Gamma_i$ comes up heads (one component of type j fails) with probability $\phi_{i,j}$ and tails (no components of type j fail) with probability $1 - \phi_{i,j}$.

We allow for a cascading failure to continue as long as there are still components operational in the system. For example, the failure of a component of type i may cause a component of type j to fail (with probability $\phi_{i,j}$), which in turn makes a component of type k fail (with probability $\phi_{j,k}$), and so on. As noted in [1], the SAVE package [3] allows for only one level of cascading, but the unlimited cascading in our model makes it significantly more difficult to analyze.

We can think of a cascading failure as a tree of simultaneously failing components. The root is the component, say of type i, whose failure *triggers* the cascade. The root's children, which are from Γ_i, are those components whose immediate failures were directly caused by the root's failure. At any non-root level of the tree, these components' failures were directly caused by the failure of their parents at the previous level. Although all the failing components in a cascade fail at the same time, we need to specify an order in which they fail for our problem to be well-defined, as we explain later. We assume the components in a tree fail in breadth-first order.

There is a single repairman who fixes failed components using a processor-sharing discipline. Specifically, if the current environment is e and there is only one failed component, which is of type i, then the repairman fixes that component at rate $\mu_{i,e}$. If there are b components currently failed, then the repairman allocates $1/b$ of his effort to each failed component, so a failed component of type i is repaired at rate $\mu_{i,e}/b$.

2 Markov Chain

We want to analyze the behavior of the system as it evolves over time. Because of the processor-sharing repair discipline and the exponential rates for the event lifetimes, it will suffice to define the state of the system as a vector containing the number of failed components of each type and the current environment. Thus, let $S = \{x = (x_1, x_2, \ldots, x_N, x_{N+1}) : 0 \leqslant x_i \leqslant r_i \ \forall i \in \Omega, \ x_{N+1} \in \mathcal{E}\}$ be the state space, and let $Z = [Z(t) : t \geqslant 0]$ be the continuous-time Markov time (CTMC) living on S keeping track of the current state of the system. (If we had instead considered a first-come-first-served repair discipline, then the state space would need to be augmented to keep track of the order in which current set of down components failed.)

We assume that Z starts in environment $0 \in \mathcal{E}$ with no components failed, i.e., state $(0, 0, \ldots, 0)$, and we now describe the CTMC's infinitesimal generator matrix $Q = (Q(x, y) : x, y \in S)$, where $Q(x, y)$ is the rate that the CTMC Z moves from state $x = (x_1, \ldots, x_N, x_{N+1})$ to state $y = (y_1, \ldots, y_N, y_{N+1})$. If $y_i = x_i$ for each $i \in \Omega$ and $y_{N+1} \neq x_{N+1}$, then (x, y) is an *environment transition* with $Q(x, y) = \nu_{x_{N+1}} \delta_{x_{N+1}, y_{N+1}}$. If $y_i = x_i - 1$ for one $i \in \Omega$, $y_j = x_j$ for each $j \in \Omega - \{i\}$, and $y_{N+1} = x_{N+1}$, then (x, y) is a *repair transition* corresponding to the repair of a component of type i, and $Q(x, y) = x_i \mu_{i, x_{N+1}} / (\sum_{j \in \Omega} x_j)$. If $y_i \geqslant x_i$ for all $i \in \Omega$ with $y_j > x_j$ for some $j \in \Omega$ and $y_{N+1} = x_{N+1}$, then (x, y) is a *failure transition* in which $y_i - x_i$ components of type i fail, $i \in \Omega$. Any other (x, y) with $x \neq y$ not falling into one of the above three categories is not possible, so $Q(x, y) = 0$. The diagonal entry $Q(x, x) = -\sum_{y \neq x} Q(x, y)$, as required for a CTMC.

We now determine the rate $Q(x, y)$ of a failure transition (x, y). First consider the case when there is no cascading failures possible, i.e., $\Gamma_i = \varnothing$ for each type i. Then the only
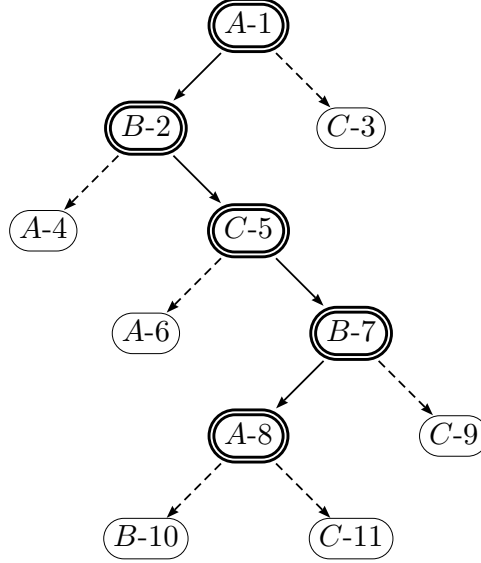
Figure 1: An example of a supertree.

possible failure transitions (x, y) have $y_i = x_i + 1$ for one $i \in \Omega$, $y_j = x_j$ for each $j \in \Omega - \{i\}$, and $y_{N+1} = x_{N+1}$, and this corresponds to a single component of type i failing. Then $Q(x, y) = (r_i - x_i)\lambda_{i,x_{N+1}}$.

Cascading failures complicate the computation of $Q(x, y)$ for a failure transition (x, y). As mentioned before, a cascading failure is modeled as a tree T built from the multiset B of simultaneously failing components, where B has $y_\ell - x_\ell$ failing components of type ℓ, $\ell \in \Omega$. A tree T has a rate $R(T)$ determined by the failure rate $(r_i - x_i)\lambda_{i,x_{N+1}}$ of the root (assumed here to be of type i), as well as the product of $\phi_{j,k}$ terms for a parent node of type j causing a child of type $k \in \Gamma_j$ to fail, and $1 - \phi_{j,k}$ terms from a node of type j *not* causing a component of type $k \in \Gamma_j$ to fail when there are components of type k up. We denote the cumulative products of $\phi_{j,k}$ terms and $1 - \phi_{j,k}$ terms by ρ and η, respectively.

$$R(T) = (r_i - x_i)\lambda_{i,x_{N+1}} \, \rho \, \eta \tag{1}$$

A difficulty arises since there can be many such trees corresponding to the multiset B of components failing in (x, y), and determining $Q(x, y)$ requires summing $R(T)$ over all possible trees T that can be constructed from B. The number of such trees grows exponentially in the number of failing components in the cascade; see [1].

Computing the rate $R(T)$ of a tree T requires that the components fail in a certain order, even though all of the failures occur simultaneously. To see why, consider the following example. Let $\Omega = \{A, B, C\}$, with $r_A = r_B = r_C = 4$. Also, let $\Gamma_A = \{B, C\}$, $\Gamma_B = \{A, C\}$, and $\Gamma_C = \{A, B\}$. Suppose that $\mathcal{E} = \{0\}$, and consider the failure transition (x, y) with $x = (2, 2, 3, 0)$ and $y = (4, 4, 4, 0)$. Thus, (x, y) corresponds to 2 components each of types A and B failing and a single component of type C failing. One possible tree T corresponding to (x, y) is shown in Figure 1.

The nodes depicted as double circles form the tree of failing components. The single circles correspond to components in some Γ_i set but did not fail. A component type j in

3

some Γ_i could have not failed because there are components of type j up at this point but its coin flip came up tails (with probability $1 - \phi_{i,j}$), or there were no more components of type j up at this point. Each node has a label of the form i-ID, where i denotes the type of the component for that node, and ID is the position of the node in a breadth-first ordering of all the nodes (single and double circles). We include the IDs just to simplify the discussion here. We call the tree of all nodes the *supertree* corresponding to the tree T of failing nodes. The supertree is used to compute $R(T)$ of T, as follows. Let u_i be the number of components of type i available in the system. Since the root is a component of type A and there are $u_A = r_A - x_A = 2$ components of type A at the start of the transition (x, y), the rate of the trigger of the cascade is $2\lambda_{A,0}$. The root then causes a component of type B to fail at node ID 2, and this occurs with probability $\phi_{A,B}$. The node at ID 3 did not fail, and at this point there are $u_C = r_C - x_C = 1 > 0$ components of type C still up, so this non-failure occurs with probability $1 - \phi_{A,C}$. Instead of stepping through the supertree one node at a time, we notice that all the $\phi_{i,j}$ terms must be included if T is to correspond to (x, y). Thus, we have $\rho = \phi_{A,B} \, \phi_{B,C} \, \phi_{C,B} \, \phi_{B,A}$. We notice the following when calculating η:

- $1 - \phi_{i,j}$ terms are included only until there are no more components of type j up.

- Once the components of type j are exhausted, all subsequent $1 - \phi_{i,j}$ terms (in breadth-first order) are not included.

- Each time we encounter a node of type j that has failed in the breadth-first enumeration of T, we decrement u_j by 1.

Keeping these observations in mind, we do the calculation for η. For component type A we have $u_A = 2$, before we traverse through T. As we do a breadth-first traversal through T, at ID 1, $u_A = 1$. We see that $u_A > 0$ until ID 8. So η includes terms $1 - \phi_{B,A}$ and $1 - \phi_{C,A}$ from IDs 2 and 5 respectively. For component type B we have $u_B = r_B - x_B = 2$ before we traverse through T in breadth-first order. Traversing through T, at ID 2, $u_B = 1$ and at ID 7, $u_B = 0$. Note that the term $1 - \phi_{A,B}$ occurs at ID 10, i.e., after ID 7. Since components of type B have already been exhausted at ID 7 we do not include $1 - \phi_{A,B}$ from ID 10. For component type C we see that $u_C = 1$ before we traverse through T and $u_C = 0$ at ID 5. Hence the only contribution we get is $1 - \phi_{A,C}$ from ID 3. Taking the product over all j, we get $\eta = (1 - \phi_{B,A}) \, (1 - \phi_{C,A}) \, (1 - \phi_{A,C})$. Therefore, $R(T)$ is $2\lambda_{A,0} \, \rho \, \eta$. In our implementation η is calculated through a data structure called *BFHist*, which we cover later.

We previously stated that the order in which the components fail in a tree must be specified for the tree's rate to be well defined. To see why, suppose instead that the components in Figure 1 failed in depth first order. The depth-first traversal of T is A-1, B-2, A-4, C-5, A-6, B-7, A-8, B-10, C-11, C-9, C-3. The number of components up of each type are $u_A = 2$, $u_B = 2$ and $u_C = 1$. In this traversal a component of type C fails at ID 5, which makes $u_C = 0$. The term $(1 - \phi_{A,C})$ occurs at IDs 11 and 3; but at either location it is not included because IDs 3 and 11 occur after ID 5. Thus, we do not get the $(1 - \phi_{A,C})$ term that we got from the breadth-first traversal; giving us a different $R(T)$. Hence, it is necessary to define the order in which the components fail, even though they fail simultaneously.

3 Algorithms

We now provide efficient algorithms for generating all possible trees and constructing the Q-matrix. A tree corresponds to a multiset of particular components failing, and cascading failures starting from different states can have the same multiset of components failing. Hence, a particular tree may correspond to several different transitions. Our algorithm generates each possible tree only once and determines all the transitions to which this tree corresponds. This avoids generating the same tree numerous times for each corresponding transition, as was originally done in [1]. Moreover, rather than building each new tree from scratch, as was done in [1], our current algorithm builds larger trees from smaller ones already considered, leading to additional savings in the overall computational effort.

Computing the rate of a given tree depends on the state from which the cascading failure began and the multiset of components that fail in the cascade. We do not actually construct the supertree in our algorithm to compute a tree's rate, but instead build a *breadth-first history* to keep track of the information necessary to compute η. We explain the process of starting a tree, building its breadth-first history and computing its rate in Algorithms 1 (SeedTrees), 2 (AddTreeLevel) and 3 (ComputeTreeRate), respectively.

SeedTrees starts the tree generation and initializes the necessary data structures for AddTreeLevel. AddTreeLevel adds a new level to an existing tree in a recursive fashion, updates ρ to include the component affected probabilities and builds the tree's breadth-first history. ComputeTreeRate computes the rate of a completed tree for all the transitions it corresponds to using ρ and η computed from the breadth-first history populated in AddTreeLevel. We will now discuss each algorithm in detail. References to line numbers in the algorithms are given within angled brackets $\langle\,\rangle$.

3.1 SeedTrees

The data structure Γ consists of $(\Gamma_i : i \in \Omega)$. We start by iterating through Ω to choose a root component with type *rootC.type* for an initial tree with one node $\langle\,1\,\rangle$. We then initialize the following data structures $\langle\,2-4\,\rangle$:

- *level*, a dynamic array to hold the failed nodes in the current bottom level of the tree (in breadth-first order);

- *nFailed*, an array that counts the number of failed components of each type in the tree; and

- *BFHist*, a data structure that is the breadth-first history of a tree.

BFHist is implemented as an array of linked list, where the array is indexed by component type. Each linked list stores the respective parents of nodes exclusive to the supertree and stores the symbol @ for the nodes that *did* fail in the tree. To determine the number of components available of a certain type, we first consider a CTMC transition (x, y) corresponding to the tree we are building, where we call x the from-state and y the to-state in the transition. Then the number of available components of a type i is computed as $r_i - n_i(x) - nFailed[i]$, where $n_i(x)$ is the number failed of type i in state x, and $nFailed[i]$ is the number failed of type i in the tree thus far in breadth-first order $\langle\,2-4\,\rangle$.

Since the root must fail, we initialize *level* with *rootC.type*. We add the symbol @ to the list at *BFHist*[*rootC.type*] to denote that one component of the root's type has failed. We update *nFailed* by setting the counter for *rootC.type* to 1 ⟨5 − 7⟩.

If the component type *rootC.type* at the root cannot cause any other components to fail, only the trivial tree of one node can be made with this type of root. We then evaluate this single-node tree's rate in ComputeTreeRate because it cannot be grown further. If $\Gamma_{rootC.type} \neq \varnothing$ (i.e., it can cause other types of components to fail) then we call AddTreeLevel to proceed with building taller trees by adding another level to the current tree ⟨9 − 13⟩.

To understand the parameters passed to subroutines ComputeTreeRate and AddTreeLevel, please refer to Equation 1.

Algorithm 1 SeedTrees(Γ)

---

where Γ is an array of ordered sets that describes which components can cause which other components to fail

1: **for** *rootC.type* ∈ Ω **do**
2:    *level* = [ ]; {Dynamic array of failed components at tree's current bottom level}
3:    *nFailed* = [0, 0, ..., 0]; {Array that counts failed components of each type in the tree}
4:    *BFHist* = [( ), ( ), ..., ( )]; {Array of linked lists that keeps a history of parent component types in breadth-first order; BFHist is indexed by component types}
5:    add *rootC.type* to *level*;
6:    *nFailed*[*rootC.type*] = 1;
7:    add @ to *BFHist*[*rootC.type*]; {Signifies one component of type *rootC.type* has failed}
8:    ρ = 1;
9:    **if** $\Gamma_{rootC.type}$ == $\varnothing$ **then**
10:       ComputeTreeRate(*nFailed*, *BFHist*, ρ, *rootC.type*);
11:    **else**
12:       AddTreeLevel(*level*, *nFailed*, *BFHist*, ρ, *rootC.type*);
13:    **end if**
14: **end for**

3.2   AddTreeLevel

Given a tree, AddTreeLevel determines all the possibilities for the next level that can be added to the current tree by taking the Cartesian product of the power sets of the nonempty Γ_i for each of the failed components at the current bottom level. For each possible next level we recursively call AddTreeLevel. We use $\mathcal{P}$ to denote the power set-like operation on an ordered set. In our realization of the power-set operation, components in each subset

maintain their relative ordering from the original set.

We implement this power set-like operation by generating all possible binary numbers with $\sum_{i=1}^{|level|} \Gamma_{level[i]}$ bits. A total of $\prod_{i=1}^{|level|} 2^{\Gamma_{level[i]}}$ such binary numbers are generated. In the binary number, 1 denotes a failed node and 0 denotes a node that is included in a Γ set but did not fail. If we have a tree where none of the leaf nodes at the current bottom level can cause any other components to fail (i.e., have empty Γ_is), we get no nextLevelPossibilities $\langle 1 \rangle$. If there are no nextLevelPossibilities, the tree cannot grow any further, so we would proceed to ComputeTreeRate.

Otherwise, we choose one possible choice for the failed components in the next level from the nextLevelPossibilities $\langle 2 \rangle$. To find out whether any new children will be added in the upcoming next level, we create a Boolean variable *addedChildFlag*, initially with the value of False $\langle 3 \rangle$. Each node in T contains three attributes: a component type i (that represents the type of the failed component of this node), an ID which is its breadth-first ID in T and a parentID which is the breadth-first ID of this node's parent in T. For each node *parentC* in the current bottom level that acts as a parent, potentially causing other nodes to fail, we iterate through all of its possible children, i.e., through $\Gamma_{parentC.type}$. If any of these children actually fail, then they will be members of the set *oneNextLevelPossibility* $\langle 4 - 6 \rangle$. Now if the redundancy of the component type we just tried to add as a child has already been exhausted, it cannot fail and hence our tree is invalid and we move on to the next possibility $\langle 7 - 9 \rangle$.

If a child of type *childC.type* has failed (i.e., it exists within nextLevelPossibilities), the relevant data structures are updated to signify this occurrence. We increment the number of failed components of *childC.type*), add the symbol @ to *BFHist[childC.type]* (to mark that a failure has occurred at this location in the tree) and ρ with $\phi_{i,j}$ $\langle 11 - 13 \rangle$. If the child does not actually fail but could have (because there are still operational components of type *childC.type* at this point), then we add *parentC.type* to *BFHist[childC.type)]* $\langle 14 - 16 \rangle$.

We do the above updates to data structures for each potential parent node in *level* and each of its children in *oneNextLevelPossibility*. If at least one child has been added, we recursively add another level to the current tree $\langle 19 - 23 \rangle$. Otherwise, this tree has not changed in this pass through AddTreeLevel and we call ComputeTreeRate to compute the rate of the finalized tree $\langle 24 - 26 \rangle$. We make sure that trees are not double counted because once a tree passes through AddTreeLevel unmodified, it is processed and discarded. We avoid making duplicate trees because each oneNextLevelPossibility is unique.

Algorithm 2 AddTreeLevel(*level, nFailed, BFHist, ρ, rootC.type*)

---

where *level* is the current level of failed components,
nFailed counts failed components by type in the tree,
BFHist is breadth-first history,
ρ is a cumulative product of component-affected probabilities,
rootC.type is the root component's type in the current tree

1: $nextLevelPossibilities = \overset{|level|}{\underset{i=1}{\times}} \mathcal{P}(\Gamma_{level[i]});$

    {Builds all possibilities for the next level (given the current level) by taking a Cartesian product of the power sets of non-empty Γ sets of failed nodes in the current level}

2: **for** $oneNextLevelPossibility \in nextLevelPossibilities$ **do**

3:     $addedAChildFlag$ = False;

4:     **for** $parentC \in level$ **do**

5:       **for** $i \in \Gamma_{parentC.type}$ **do**

6:         **if** $\exists\ childC \in oneNextLevelPossibility : childC.type == i$ **and** $childC.parentID == parentC.ID$ **then**

7:           $addedAChildFlag$ = True;

8:           **if** $nFailed[childC.type] < r_{childC.type}$ **then**

9:             $nFailed[childC.type] = nFailed[childC.type] + 1;$

10:           **else**

11:             goto line 2; {Invalid Tree, it requires more components than available}

12:           **end if**

13:           add @ to $BFHist[childC.type]$; {One component of type $childC.type$ has failed}

14:           $\rho = \rho\ \phi_{parentC.type,\ childC.type};$
          {Update rate with appropriate component-affected probability}

15:         **else**

16:           add $parentC.type$ to $BFHist[childC.type]$; {One component of type $childC.type$ has not failed, but was present in $\Gamma_{parentC.type}$}

17:         **end if**

18:       **end for**

19:     **end for**

20:     **if** $addedAChildFlag$ **then**

21:       AddTreeLevel($oneNextLevelPossibility$, $nFailed$, $BFHist$, ρ, $rootC.type$);
      {Tree can be grown further}

22:     **else**

23:       ComputeTreeRate($nFailed$, $BFHist$, ρ, $rootC.type$);
      {Current tree is complete because it cannot be grown further}

24:     **end if**

25: **end for**

3.3   ComputeTreeRate

In ComputeTreeRate, for a given tree, we determine all of the transitions (x, y) in the Q-matrix that use this tree, and then update the total rates of each of those transitions (x, y) by adding in the rate of the current tree to the current rate for (x, y). However, even if several transitions use the same tree, the rate of the tree may differ for those transitions, since the failure rate of the root and η are transition dependent.

    Let $x' = (x_1, x_2, \dots, x_N)$ and $S' = (x_1, x_2, \dots, x_N) : 0 \leqslant x_i \leqslant r_i, \forall i \in \Omega$. In line 1 of ComputeTreeRate, we only loop through x' and not x because the structure of a tree does

not depend on the environment; only the tree's rate does. For each component type i in Ω, we calculate the number of components that are available in the system, $u_i = r_i - n_i(x)$. The number u_i determines until which point, while traversing through the $BFHist[i]$, we can still have components of type i that could have failed but did not fail $\langle 4 \rangle$. We cannot have components of a type that could have failed but did not after a point in the tree where its type has been exhausted. A type is exhausted when the total number of failed components of the type (i.e., the sum of the components already failed in x' and in the tree in breadth-first order until the current node in the $BFHist[i]$) equals the redundancy of the type. Each time we encounter the symbol @ in the $BFHist$, we reduce the number available by one $\langle 5 - 7 \rangle$. If there are components still available, we update η $\langle 8 - 9 \rangle$. As soon as u_i reaches zero, we break because that type has then been exhausted, i.e., there are no more components that could have failed for the type $\langle 10 - 12 \rangle$.

Once we have finished calculating η, we loop through all environments since the same tree can be used for transitions occurring in different environments $\langle 15 \rangle$. We convert x' into a state x by appending an environment e $\langle 16 \rangle$. We generate a state y by adding *nFailed* to x' and appending environment e $\langle 17 \rangle$. Invalid states y will be generated in some instances because simply adding *nFailed* will cause some component types' number failed to exceed their redundancies $\langle 18 - 20 \rangle$. The failure rate of the root is calculated as the number of components of type *rootC.type* up in the system, multiplied by the failure rate of the component in the environment of x $\langle 21 \rangle$. With all parts of the rate calculation done, we add the rate of the current tree to the current cumulative rate for transition (x, y) in the Q-matrix $\langle 22 \rangle$.

Algorithm 3 ComputeTreeRate(*nFailed*, *BFHist*, ρ, *rootC.type*)

---

where *level* is the current bottom level of failed components,
nFailed counts the failed components of each type in the tree,
BFHist is breadth-first history,
ρ is a cumulative product of component-affected probabilities,
rootC.type is the root's type in the tree

1: **for** $x' \in S'$ **do**
2:    $\eta = 1$; {Cumulative product of complement probabilities of components that could have failed but did not}
3:    **for** $i \in \Omega$ **do**
4:       $u_i = r_i - x_i'$;
5:       **for** *parentC.type* in $BFHist[i]$ **do**
6:          **if** *parentC.type* == @ **then**
7:             $u_i = u_i - 1$;
8:          **else if** $u_i > 0$ **then**
9:             $\eta = \eta \ (1 - \phi_{parentC.type,\ i})$;
10:          **else**
11:             break; {u_i must be greater than 0 or else there cannot be any more failed nodes of type i}
12:          **end if**

```
13:       end for
14:     end for
15:     for $e \in \mathcal{E}$ do
16:         $x = (x', e)$;
17:         $y = (x' + nFailed, e)$;
18:         if $y$ is not a valid state then
19:             continue;
20:         end if
21:         $rootFailureRate = (r_{rootC.type} - x[rootC.type]) * \lambda_{rootC.type,\, x_{N+1}}$;
22:         $Q(x, y) = Q(x, y) + rootFailureRate * \rho * \eta$;
23:     end for
24: end for
```

4 Experimental

4.1 Nested Experimental

References

[1] S. M. Iyer, M. K. Nakayama, and A. V. Gerbessiotis, "A Markovian dependability model with cascading failures," *IEEE Transactions on Computers*, vol. 139, pp. 1238–1249, 2009.

[2] M. Xie, Y. S. Dai, and K. Poh, *Computing Systems Reliability: Models and Analysis.* New York: Kluwer Academic, 2004.

[3] A. Blum, P. Heidelberger, S. S. Lavenberg, M. K. Nakayama, and P. Shahabuddin, "Modeling and analysis of system availability using SAVE," in *Proceedings of the 23rd International Symposium on Fault Tolerant Computing*, pp. 137–141, 1994.