

Working Title

M. Sanghavi, S. Tadepalli, M. Nakayama

January 2013

1 Model

We work with the stochastic model of [1], which considers the evolution over time of a repairable dependability system operating in a randomly changing environment. The system consists of a collection $\Omega = \{1, 2, \dots, N\}$ of $N < \infty$ component types. Each component type $i \in \Omega$ has a redundancy $1 \leq r_i < \infty$, and the r_i components of type i are assumed to be identical. A component can be either operational (up) or failed (down).

The environment changes randomly within a set $\mathcal{E} = \{0, 1, 2, \dots, L\}$. For example, we can think of the environment as representing the current load on the system, and if there are two possible environments, 0 and 1, then 0 (resp., 1) may represent a low (resp., high) load. Once the environment enters $e \in \mathcal{E}$, it remains there for an exponentially distributed amount of time with rate $\nu_e > 0$, after which the environment changes to e' with probability $\delta_{e,e'} \geq 0$, where $\delta_{e,e} = 0$ and $\sum_{e' \in \mathcal{E}} \delta_{e,e'} = 1$. We assume the matrix $\delta = (\delta_{e,e'} : e, e' \in \mathcal{E})$ is irreducible; i.e., for each $e, e' \in \mathcal{E}$, there exists $k \geq 1$ and a sequence $e_0 = e, e_1, e_2, \dots, e_k = e'$ with each $e_i \in \mathcal{E}$ such that $\prod_{i=0}^{k-1} \delta_{e_i, e_{i+1}} > 0$. In other words, it is possible to eventually move from environment e to environment e' .

The components in the system can randomly fail and then be repaired. When the environment is $e \in \mathcal{E}$, the failure rate and repair rate of each component of type i are $\lambda_{i,e} > 0$ and $\mu_{i,e} > 0$, respectively. If there is only one environment e , i.e., $|\mathcal{E}| = 1$, then the lifetimes and repair times of components of type i are exponentially distributed with rates $\lambda_{i,e}$ and $\mu_{i,e}$, respectively. Exponential distributions are frequently used to model lifetimes of hardware and software components; e.g., see [2]. We assume that all operating components of a type i have the same failure rate $\lambda_{i,e}$ in environment e . Thus, in a system with redundancies for which not all components of a type are needed for operation of the system, the extras are “hot spares” since they fail at the same rate as the main components.

Our model includes probabilistic instantaneous cascading failures, which occur as follows. The ordered set Γ_i specifies the types of components that a failure of a type- i component can cause to simultaneously fail. When a component of type i fails, it causes a single component of type $j \in \Gamma_i$ to fail at the same time with probability $\phi_{i,j}$ (if there are components of type j up). The events that the individual components types $j \in \Gamma_i$ fail immediately are independent. Thus, when a component of type i fails, there are independent “coin flips” to determine which components in Γ_i fail, where the coin flip for $j \in \Gamma_i$ comes up heads (one component of type j fails) with probability $\phi_{i,j}$ and tails (no components of type j fail) with probability $1 - \phi_{i,j}$.

We allow for a cascading failure to continue as long as there are still components operational in the system. For example, the failure of a component of type i may cause a component of type j to fail (with probability $\phi_{i,j}$), which in turn causes a component of type k to fail (with probability $\phi_{j,k}$), and so on. As noted in [1], the SAVE package [3] allows for only one level of cascading, but the unlimited cascading in our model makes it significantly more difficult to analyze.

We can think of a cascading failure as a tree of simultaneously failing components. The root is the component, say of type i , whose failure *triggers* the cascade. The root's children, which are from Γ_i , are those components whose immediate failures were directly caused by the root's failure. At any non-root level of the tree, these components' failures were directly caused by the failure of their parents at the previous level. Although all the failing components in a cascade fail at the same time, we need to specify an order in which they fail for our problem to be well-defined, as we explain later. We assume the components in a tree fail in breadth-first order.

There is a single repairman who fixes failed components using a processor-sharing discipline. Specifically, if the current environment is e and there is only one failed component, which is of type i , then the repairman fixes that component at rate $\mu_{i,e}$. If there are b components currently failed, then the repairman allocates $1/b$ of his effort to each failed component, so a failed component of type i is repaired at rate $\mu_{i,e}/b$.

2 Markov Chain

We want to analyze the behavior of the system as it evolves over time. Because of the processor-sharing repair discipline and the exponential distributions for the event lifetimes, it will suffice to define the state of the system as a vector containing the number of failed components of each type and the current environment. Thus, let $S = \{x = (x_1, x_2, \dots, x_N, x_{N+1}) : 0 \leq x_i \leq r_i \ \forall i \in \Omega, x_{N+1} \in \mathcal{E}\}$ be the state space, and let $Z = [Z(t) : t \geq 0]$ be the continuous-time Markov time (CTMC) living on S keeping track of the current state of the system. (If we had instead considered a first-come-first-served repair discipline, then the state space would need to be augmented to keep track of the order in which current set of down components failed.)

We assume that Z starts in environment $0 \in \mathcal{E}$ with no components failed, i.e., state $(0, 0, \dots, 0)$, and we now describe the CTMC's infinitesimal generator matrix $Q = (Q(x, y) : x, y \in S)$, where $Q(x, y)$ is the rate that the CTMC Z moves from state $x = (x_1, \dots, x_N, x_{N+1})$ to state $y = (y_1, \dots, y_N, y_{N+1})$. If $y_i = x_i$ for each $i \in \Omega$ and $y_{N+1} \neq x_{N+1}$, then (x, y) is an *environment transition* with $Q(x, y) = \nu_{x_{N+1}} \delta_{x_{N+1}, y_{N+1}}$. If $y_i = x_i - 1$ for one $i \in \Omega$, $y_j = x_j$ for each $j \in \Omega - \{i\}$, and $y_{N+1} = x_{N+1}$, then (x, y) is a *repair transition* corresponding to the repair of a component of type i , and $Q(x, y) = x_i \mu_{i, x_{N+1}} / (\sum_{j \in \Omega} x_j)$. If $y_i \geq x_i$ for all $i \in \Omega$ with $y_j > x_j$ for some $j \in \Omega$ and $y_{N+1} = x_{N+1}$, then (x, y) is a *failure transition* in which $y_i - x_i$ components of type i fail, $i \in \Omega$. Any other (x, y) with $x \neq y$ not falling into one of the above three categories is not possible, so $Q(x, y) = 0$. The diagonal entry $Q(x, x) = -\sum_{y \neq x} Q(x, y)$, as required for a CTMC.

We now determine the rate $Q(x, y)$ of a failure transition (x, y) . First consider the case when there is no cascading failures possible, i.e., $\Gamma_i = \emptyset$ for each type i . Then the only

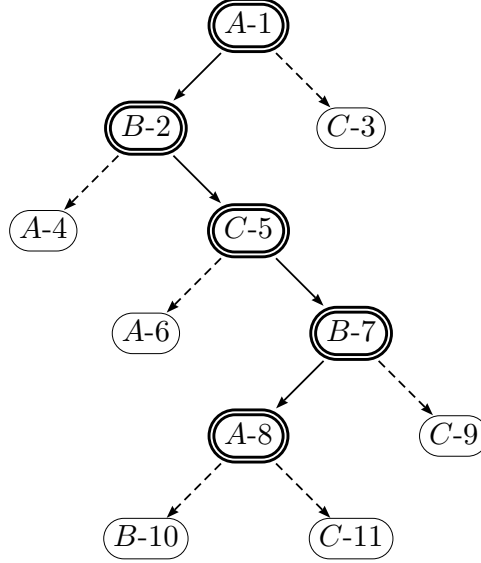


Figure 1: An example of a supertree.

possible failure transitions (x, y) have $y_i = x_i + 1$ for one $i \in \Omega$, $y_j = x_j$ for each $j \in \Omega - \{i\}$, and $y_{N+1} = x_{N+1}$, and this corresponds to a single component of type i failing. Then $Q(x, y) = (r_i - x_i)\lambda_{i, x_{N+1}}$.

Cascading failures complicate the computation of $Q(x, y)$ for a failure transition (x, y) . As mentioned before, a cascading failure is modeled as a tree T built from the multiset B of simultaneously failing components, where B has $y_\ell - x_\ell$ failing components of type ℓ , $\ell \in \Omega$. A tree T has a rate $R(T)$ determined by the failure rate $(r_i - x_i)\lambda_{i, x_{N+1}}$ of the root (assumed here to be of type i), as well as the product of $\phi_{j,k}$ for a parent node of type j causing a child of type $k \in \Gamma_j$ to fail, and $1 - \phi_{j,k}$ from a node of type j *not* causing a component of type $k \in \Gamma_j$ to fail when there are components of type k up. A difficulty arises since there can be many such trees corresponding to the multiset B of components failing in (x, y) , and determining $Q(x, y)$ requires summing $R(T)$ over all possible trees T that can be constructed from B . The number of such trees grows exponentially in the number of failing components in the cascade; see [1].

[The following is not complete and will need to be changed for another example.] Computing the rate $R(T)$ of a tree T requires that the components fail in a certain order, even though all of the failures occur simultaneously. To see why, consider the following example. Let $\Omega = \{A, B, C\}$, with $r_A = r_B = r_C = 4$. Also, let $\Gamma_A = \{B, C\}$, $\Gamma_B = \{A, C\}$, and $\Gamma_C = \{A, B\}$. Suppose that $\mathcal{E} = \{0\}$, and consider the failure transition (x, y) with $x = (2, 2, 2, 0)$ and $y = (4, 4, 3, 0)$. Thus, (x, y) corresponds to 2 components each of types A and B failing and a single component of type C failing. One possible tree T corresponding to (x, y) is shown in Figure 1.

The nodes depicted as double circles form the tree of failing components. The single circles correspond to components in some Γ_i set but did not fail. A component type j in some Γ_j could have not failed because there are components of type j up at this point but

its coin flip came up tails (with probability $1 - \phi_{i,j}$), or there were no more components of type j up at this point. Each node has a label of the form $t-i$, where t denotes the type of the component for that node, and i is the ID, which is the position of the node in a breadth-first ordering of all the nodes (single and double circles). We include the IDs just to simplify the discussion here. We call the tree of all nodes the *supertree* corresponding to the tree T of failing nodes. The supertree is used to compute the rate $R(T)$ of T , as follows. Since the root is a component of type A and there are $r_A - x_A = 2$ components of type A at the start of the transition (x, y) , the rate of the trigger of the cascade is $2\lambda_{A,0}$. The root then causes a component of type B to fail at node ID 2, and this occurs with probability $\phi_{A,B}$. The node at ID 3 did not fail, and at this point there are $r_C - x_C = 2 > 0$ components of type C still up, so this nonfailure occurs with probability $1 - \phi_{A,C}$.

3 Algorithms

We now provide efficient algorithms for generating all possible trees and constructing the Q -matrix. A tree corresponds to a set of particular components failing, and cascading failures starting from different states can have the same set of components failing. Hence, a particular tree may correspond to several different transitions. Our algorithm generates each possible tree only once and determines all the transitions to which this tree corresponds. This avoids generating the same tree numerous times for each corresponding transition, as was originally done in [1]. Because the number of trees grows exponentially in the number of components in the model, our current algorithm significantly reduces the total computational effort. Moreover, rather than building each new tree from scratch, as was done in [1], our current algorithm builds larger trees from smaller ones already considered, leading to additional savings in the computation.

Computing the rate of a given tree depends on the state from which the cascading failure began and the set of components that fail in the cascade. The failure of a component type i can probabilistically cause each component type from a set Γ_i to probabilistically simultaneously fail when i fails; i.e., not all components in Γ_i actually fail in a given cascade. We also consider the supertree from a given tree by adding in the components that did not fail from the Γ_i that were used, which is necessary to compute the rate of a tree. We do not actually construct the supertree in our algorithm, but instead build a *breadth-first history* to keep track of the information necessary to compute a tree's rate. A breadth-first history accounts for the contribution to the tree's rate of the components that could have failed (i.e., belong to the set Γ_i of a node of type i that actually did fail), but did not. We explain the process of building a breadth-first history and using it to compute the rate of a tree in Algorithms 2 (AddTreeLevel) and 3 (ComputeTreeRate), respectively.

SeedTrees (Algorithm 1), starts the tree generation and initializes the necessary data structures for AddTreeLevel. AddTreeLevel adds a new level to an existing tree in a recursive fashion, updates the cumulative failed probability for the tree for the components that actually failed in the cascade as well as builds the tree's breadth-first history. ComputeTreeRate computes the rate of a completed tree for all the transitions it corresponds to using the cumulative failed probability and the breadth-first history populated in AddTreeLevel. We will now discuss each algorithm in detail. Line numbers from each the algorithms are

given in their corresponding texts within angled brackets $\langle \rangle$.

3.1 SeedTrees

The data structure $\Gamma = (\Gamma_i : i \in \Omega)$ describes how cascading failures can occur, with Γ_i as the (ordered) set of components that probabilistically and instantaneously fail when a component of type i fails. We start by iterating through Ω , the set of all component types, to choose a root component, $rootC$, for an initial tree with one node $\langle 1 \rangle$. We then initialize the following data structures:

- *level*, a dynamic array to hold the failed nodes in the current bottom level of the tree (in breadth-first order);
- *nFailed*, a list that counts the number of failed components of each type in the tree; and
- *BFHist*, a data structure that is the breadth-first history of a tree.

BFHist is implemented as an array of linked lists indexed by component type. Each linked list stores the respective parents of nodes exclusive to the supertree (i.e., each node that did not fail but could have if its component type's redundancy was not exhausted) and stores the symbol @ for the nodes that *did* fail in the tree. *BFHist* plays the role of a supertree in determining whether to include the complement probabilities of nodes that did not fail. A complement probability is included for a node that did not fail only if there are still components of its type available in the system at that point. To determine the number of components available of a certain type, we first consider a CTMC transition (x, y) corresponding to the tree we are building, where we call x the from-state and y the to-state in the transition. Then the number of available components of a type i is computed as $r_i - n_i(x) - nFailed[i]$, where $n_i(x)$ is the number failed of type i in state x , and $nFailed[i]$ is the number failed of type i in the tree thus far in breadth-first order $\langle 2 - 4 \rangle$.

Since the root must fail, we initialize *level* with $rootC$. We add the symbol @ to $BFHist[rootC]$ to denote that one component the root's type has failed. We update *nFailed* by setting the counter for $rootC$'s type to 1 $\langle 5 - 7 \rangle$.

If the component $rootC$ at the root cannot cause any other components to fail, only the trivial tree of one node can be made with this type of root. We then evaluate this single-node tree's rate in `ComputeTreeRate` because it cannot be grown further. If $\Gamma_{rootC} \neq \emptyset$, i.e., it can cause other types of components to fail, then we call `AddTreeLevel` to proceed with building taller trees by adding another level to the current tree $\langle 8 - 13 \rangle$.

In ρ we only keep track of the cumulative product of the component-affected probabilities $\phi_{i,j}$ for the transition rate of the tree, which entails multiplying $\phi_{i,j}$ for all edges from a parent of type i to a child of type j in the tree. The total transition rate of a completed tree consists of three parts: the failure rate of the root, cumulative production of the $\phi_{i,j}$'s computed in `AddTreeLevel`, which is ρ , and the cumulative product of the complement probabilities of components that could have failed but did not, which is η . The failure rate of the root for a transition (x, y) corresponding to this tree is given by the $n_{rootC} \lambda_{rootC, e}$ term, where $\lambda_{(rootC), e}$ is the failure rate of $rootC$ in environment e . Each component type has an

environment dependent failure rate. The cumulative product of complement probabilities consists of product terms of the form $1 - \phi_{i,j}$ for all edges from parent of type i to a child of type j in the supertree, where failures did not occur but could have because of components of type j were still up. We multiply the $1 - \phi_{i,j}$ terms and the $n_{rootC} \lambda_{rootC,e}$ term later in `ComputeTreeRate` as they depend on the transition (x, y) . Since the tree initially has only one node, there are no failed nodes from a cascade (i.e., nodes have been caused to fail by other nodes). Hence, we pass 1 as the value for ρ to the subroutines `AddTreeLevel` and `ComputeTreeRate`.

Algorithm 1 `SeedTrees(Γ)`

where Γ is an array of ordered sets that describes which components can cause which other components to fail

```

1: for  $rootC \in \Omega$  do
2:    $level = [ ]$ ; {dynamic array of failed components at tree's current bottom level}
3:    $nFailed = (0, 0, \dots, 0)$ ; {counts failed components of each type in the tree}
4:    $BFHist = [(), (), \dots, ()]$ ; {an array of linked lists that keeps a breadth-first history of
   trees; array is indexed by component type; linked list for each component type stores
   parents in breadth-first order}
5:   add  $rootC$  to  $level$ ;
6:    $nFailed[rootC] = 1$ ;
7:   add @ to  $BFHist[rootC]$ ; {signifies one component of type  $rootC$  has failed}
8:   if  $\Gamma_{rootC} == \emptyset$  then
9:     ComputeTreeRate( $nFailed, BFHist, 1, rootC$ );
10:  else
11:    AddTreeLevel( $level, nFailed, BFHist, 1, rootC$ );
12:  end if
13: end for

```

3.2 AddTreeLevel

In `AddTreeLevel`, given a tree, we determine all the possibilities for the next level that can be added to the current tree by taking the Cartesian product of the power sets of the Γ_i for each of the failed components at the current bottom level. For each next level we recursively call `AddTreeLevel`. We use \mathcal{P} to denote the power set-like operation on an ordered set. In our realization of the power-set operation, components in each subset maintain their relative ordering from the original set.

We implement this power set-like operation by generating all possible binary numbers with $\sum_{i=1}^{|level|} \Gamma_{level[i]}$ bits. A total of $\prod_{i=1}^{|level|} 2^{\Gamma_{level[i]}}$ such binary numbers are generated. In the

binary number, 1 denotes a failed node, 0 denotes a node that is included in a Γ set but did not fail. If it so happens that we have a tree where none of the leaf nodes at the current bottom level can cause any other components to fail (i.e., have empty Γ_i s), we get no *nextLevelPossibilities* $\langle 1 \rangle$. If there are no *nextLevelPossibilities*, we immediately proceed to *ComputeTreeRate* $\langle 2 - 4 \rangle$.

Otherwise, we choose one possible choice for the failed components in the next level to work with from the *nextLevelPossibilities* $\langle 5 \rangle$. To find out whether any new children will be added in the upcoming next level, we create a Boolean variable *addedChildFlag*, initially with the value of False $\langle 6 \rangle$. For each node *parentC* in the current bottom level that acts as a parent, potentially causing other nodes to fail, we iterate through all of its possible children, i.e., through $\Gamma_{parentC}$. If any of these children actually fail, then they will be members of the set *oneNextLevelPossibility* $\langle 7 - 9 \rangle$. Now if the redundancy of the component type we just tried to add as a child has already been exhausted, it cannot fail and hence our tree is invalid and we move on to the next possibility $\langle 10 - 12 \rangle$.

If it is indeed possible to add a child of *Type(childC)* we flag this occurrence and update the relevant data structures. We increment the number of failed components of *Type(childC)*. We add the symbol @ to *BFHist[ChildC]* to mark that a failure has occurred at this location in the tree. We update tree rate with the component-affected probability of the parent causing the child to fail $\langle 13 - 16 \rangle$. If the child does not actually fail but could have (because there are still operational components of the child's type at this point), then we add *parentC* to *BFHist[ChildC]*. *BFHist* comes in use later in *ComputeTreeRate*, to determine when to multiply the current tree rate with the $1 - \phi_{i,j}$ terms for the components that could have failed by did not $\langle 17 - 19 \rangle$.

We do the above updates to data structures for each potential parent node in *level* and each of its children in *oneNextLevelPossibility* $\langle 7 - 21 \rangle$. If at least one child has been added, we recursively add another level to the current tree $\langle 22 - 23 \rangle$. Otherwise, this tree has not changed in this pass through *AddTreeLevel* and we call *ComputeTreeRate* to compute the rate of the finalized tree $\langle 24 - 26 \rangle$. We make sure that trees are not double counted because once a tree passes through *AddTreeLevel* unmodified, it is processed and discarded. We do not make duplicate trees because each *oneNextLevelPossibility* is unique.

Algorithm 2 *AddTreeLevel*(*level*, *nFailed*, *BFHist*, ρ , *rootC*)

where *level* describes failed components,

nFailed counts failed components by type,

BFHist is breadth-first history,

ρ is the cumulative probability of components that failed,

rootC is the root component of the current tree

- 1: $nextLevelPossibilities = \prod_{i=1}^{|level|} \mathcal{P}(\Gamma_{level[i]});$
 {Builds set of all possible nodes in next level as Cartesian product of power sets of Γ s}
- 2: **if** *Empty*(*nextLevelPossibilities*) **then**
- 3: *ComputeTreeRate*(*nFailed*, *BFHist*, ρ , *rootC*);
 {current tree cannot be grown further because its leaf nodes at the current bottom

```

    level have empty  $\Gamma$ 
4: end if
5: for  $oneNextLevelPossibility \in nextLevelPossibilities$  do
6:    $addedChildFlag = \text{False};$ 
7:   for  $parentC \in level$  do
8:     for  $childC \in \Gamma_{parentC}$  do
9:       if  $childC \in oneNextLevelPossibility$  then
10:        if  $nFailed[childC] == r_{childC}$  then
11:          goto line 3; {invalid tree, requires more components than available in system}
12:        end if
13:         $addedChildFlag = \text{True};$ 
14:         $nFailed[childC] = nFailed[childC] + 1;$ 
15:        add @ to  $BFHist[childC]$ ; {signifies one component of type childC has failed}
16:         $\rho = \rho * \phi_{parentC, childC};$ 
        {update rate with component-affected probability}
17:      else
18:        add  $parentC$  to  $BFHist[childC]$ ; {signifies one component of type childC has
        not failed, but was present in  $\Gamma_{parentC}$ }
19:      end if
20:    end for
21:  end for
22:  if  $addedChildFlag$  then
23:    AddTreeLevel( $oneNextLevelPossibility, nFailed, BFHist, \rho, rootC$ );
    {tree can be grown further}
24:  else
25:    ComputeTreeRate( $nFailed, BFHist, \rho, rootC$ );
    {current tree is completed because it cannot be grown further}
26:  end if
27: end for

```

3.3 ComputeTreeRate

In ComputeTreeRate, for a given tree, we determine all of the transitions (x, y) in the Q -matrix that use this tree, and then update the total rates of each of those transitions (x, y) by adding in the rate of the current tree to the current rate for (x, y) . However, even if several transitions use the same tree, the rate of the tree may differ for those transitions. Thus, we need to compute a separate tree rate for each transition since the failure rate of the root and cumulative product of complement probabilities are transition dependent.

Let state $x = (x_1, x_2, \dots, x_{|\Omega|}, e)$ and let $x' = (x_1, x_2, \dots, x_{|\Omega|})$. Let the state space $S = (x_1, x_2, \dots, x_{|\Omega|}, e)$ and let $S' = (x_1, x_2, \dots, x_{|\Omega|}) : 0 \leq x_i \leq r_i \forall i \in \Omega$. We only loop through x' and not x because the structure of the trees do not depend on the environment; only the rate does $\langle 1 \rangle$. The variable η denotes the product of the complement probabilities of all the nodes in the tree that did not fail, but could have if their respective type's redundancy

had not been exhausted $\langle 2 \rangle$. For each component type in Ω , we calculate the number of components that are available in the system as the type's redundancy minus the number already failed in the from-state. The number *compsAvailable* determines until which point, while traversing through the *BFHist*, we can still have components of a given type that could have failed but did not fail $\langle 4 \rangle$. We cannot have components that *could have failed but did not* after a point in the tree where their type has been exhausted. A type is exhausted when the total number of failed components of the type in the system, i.e., the sum of the components failed in the from-state and in the tree in breadth-first order until the current node, equals the redundancy of the type. Each time we encounter the symbol @ in the *BFHist*, we reduce the number available by one $\langle 5 - 7 \rangle$. If there are components still available, we update η $\langle 8 - 9 \rangle$. As soon as *compsAvailable* reaches zero, we break because that type has then been exhausted i.e., there are no more components that could have failed for the type $\langle 10 - 12 \rangle$.

Once we have finished calculating the η , we loop through all environments since the same tree can be used for transitions occurring in different environments $\langle 15 \rangle$. We generate from-states by adding an environment e to x' $\langle 16 \rangle$. We generate to-states y (with an environment e) by adding *nFailed* to x' $\langle 17 \rangle$. Invalid to-states will be generated in some instances, because simply adding *nFailed* will cause component types' number failed to exceed their redundancies $\langle 18 - 20 \rangle$. The failure rate of the root is calculated as the number of components of the *rootC*'s type up in the system, multiplied by the failure rate of the component in the environment of y $\langle 21 \rangle$. With all parts of the rate calculation done, we add the rate of the current tree to the current cumulative rate for transition(x, y) in the *Q*-matrix $\langle 23 \rangle$.

Algorithm 3 ComputeTreeRate(*nFailed*, *BFHist*, ρ , *rootC*)

where *level* describes failed components,
nFailed counts failed components of each type in the tree,
BFHist is breadth-first history,
 ρ is a cumulative probability of components that failed,
rootC is the root component of the current tree

```

1: for  $x' \in S'$  do
2:    $\eta = 1$ ; {cumulative product of complement probabilities of components that could
   have failed but did not}
3:   for  $comp \in \Omega$  do
4:      $compsAvailable = r_{comp} - x'[comp]$ ;
5:     for  $parentC \in BFHist[comp]$  do
6:       if  $parentC == @$  then
7:          $compsAvailable = compsAvailable - 1$ ;
8:       else if  $compsAvailable > 0$  then
9:          $\eta = \eta * (1 - \phi_{parentC, comp})$ ;
10:      else
11:        break; {compsAvailable must equal 0 which means there cannot be any more
        nodes that could have failed of Type(comp)}

```

```

12:     end if
13:   end for
14: end for
15: for  $e \in \mathcal{E}$  do
16:    $x = (x', e)$ ;
17:    $y = (x' + nFailed, e)$ ;
18:   if  $y$  is not a valid state then
19:     continue;
20:   end if
21:    $rootFailureRate = (r_{rootC} - x[rootC]) * \lambda_{rootC, e}$ ;
22: end for
23:  $Q(x, y) = Q(x, y) + rootFailureRate\eta$ ;
24: end for

```

4 Experimental

4.1 Nested Experimental

References

- [1] S. M. Iyer, M. K. Nakayama, and A. V. Gerbessiotis, “A Markovian dependability model with cascading failures,” *IEEE Transactions on Computers*, vol. 139, pp. 1238–1249, 2009.
- [2] M. Xie, Y. S. Dai, and K. Poh, *Computing Systems Reliability: Models and Analysis*. New York: Kluwer Academic, 2004.
- [3] A. Blum, P. Heidelberger, S. S. Lavenberg, M. K. Nakayama, and P. Shahabuddin, “Modeling and analysis of system availability using SAVE,” in *Proceedings of the 23rd International Symposium on Fault Tolerant Computing*, pp. 137–141, 1994.